

# Linux\_Process\_Management\_Guide

---

## Linux Process Management Guide

---

**Author:** Divyanshi Thapa

**Date:** September 26, 2025

---

### ps accepts options in various formats:

- not preceded with a dash (BSD style)
- preceded with a dash (UNIX style)
- preceded with two dashes (GNU style)

## Table of Contents

1. [Viewing All Processes](#)
2. [Simple filtering](#)
3. [Listing threads](#)
4. [Listing Child processes](#)
5. [Controlling the Output](#)
6. [Process Tree](#)
7. [Real-Time Process Monitoring](#)
8. [Adjusting Process Priority](#)
9. [CPU Affinity \(Bind Process to CPU Core\)](#)
10. [I/O Scheduling Priority](#)
11. [File Descriptors Used by a Process](#)
12. [Trace System Calls of a Process](#)
13. [Find Process Using a Port](#)
14. [Per-Process Statistics](#)
15. [Control Groups for Resource Limits](#)
16. [Alternatives to nice / renice](#)
17. [Cheat Sheet](#)

---

NOTE: To find a process by its name like "sleep 3000": `sudo pgrep -f "sleep 1000"`

NOTE: To kill a process if you know its name like "sleep 300" : `sudo pkill -f "sleep 300"`

## 1. List Processes

**Command:**

```
ps aux
```

**Explanation:**

- **ps** → process status
- **a** → show processes for all users
- **u** → show user/owner of process
- **x** → includes off terminal commands

**Example Output:**

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.1	167500	1100	?	Ss	Sep25	0:05	/sbin/init
vibhu	1234	1.2	1.5	274532	15632	?	S1	10:15	0:12	/usr/bin/python3
				script.py						
mysql	2001	0.5	2.0	450000	20988	?	Ssl	Sep25	1:02	/usr/sbin/mysqld

**Sample code:**

```
vboxuser@ubuntu:~$ ps au
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START  TIME COMMAND
vboxuser  1932  0.0  0.0 235668  5876  tty2      Ssl+ 04:42  0:00 /usr/libexec/gdm-wayland-session env GNOME_SHELL_SES
vboxuser  1944  0.0  0.2 298236 16316  tty2      Sl+  04:42  0:00 /usr/libexec/gnome-session-binary --session=ubuntu
vboxuser  5027  0.0  0.0 11260   5544 pts/1      Ss   07:45  0:01 bash
vboxuser  8660  600 0.0 13616   4576 pts/1      R+   11:22  0:00 ps au
vboxuser@ubuntu:~$ ps u
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START  TIME COMMAND
vboxuser  1932  0.0  0.0 235668  5876  tty2      Ssl+ 04:42  0:00 /usr/libexec/gdm-wayland-session env GNOME_SHELL_SES
vboxuser  1944  0.0  0.2 298236 16316  tty2      Sl+  04:42  0:00 /usr/libexec/gnome-session-binary --session=ubuntu
vboxuser  5027  0.0  0.0 11260   5544 pts/1      Ss   07:45  0:01 bash
vboxuser  8661  100 0.0 13616   4592 pts/1      R+   11:22  0:00 ps u
vboxuser@ubuntu:~$ ps a
  PID TTY      STAT   TIME COMMAND
 1932  tty2      Ssl+  0:00 /usr/libexec/gdm-wayland-session env GNOME_SHELL_SESSION_MODE=ubuntu /usr/bin/gnome-sessio
 1944  tty2      Sl+   0:00 /usr/libexec/gnome-session-binary --session=ubuntu
 5027  pts/1      Ss    0:01 bash
 8662  pts/1      R+   0:00 ps a
vboxuser@ubuntu:~$ ps x
  PID TTY      STAT   TIME COMMAND
 1860  ?        Ss    0:12 /usr/lib/systemd/systemd --user
 1861  ?        S     0:00 (idle)
```

**Command :** `ps -e`

Note: see a more detailed output by using the **-f** in last.

Note : **-f** adds on the additional info in all the other commands.

## 2. Simple filtering

**Command:** `ps -C processname`

searching for a particular process by name with the **-C** option.

**Command:** `ps -p processid`

filter based on a list of process ids using the -p flag.

**Command:** `ps -u username`

search by the user name by specifying it in the -u option.

**Sample code:**

```
vboxuser@ubuntu:~$ ps -p 8707
  PID TTY          TIME CMD
  8707 ?        00:00:00 kworker/3:0-mm_percpu_wq
vboxuser@ubuntu:~$ ps -u vbosuser | head -5
error: user name does not exist
```

**Usage:**

```
ps [options]
```

```
Try 'ps --help <simple|list|output|threads|misc|all>'
or 'ps --help <s|l|o|t|m|a>'
for additional help text.
```

For more details see `ps(1)`.

```
vboxuser@ubuntu:~$ ps -u vboxuser | head -5
  PID TTY          TIME CMD
 1860 ?        00:00:12 systemd
 1864 ?        00:00:00 (sd-pam)
 1875 ?        00:00:17 pipewire
 1876 ?        00:00:00 pipewire
vboxuser@ubuntu:~$
```

### 3. Listing Threads

**Command:** `ps -C gedit -L`

- Todo: search a process named gedit and list all its threads.

**Note :** you can ask for a more informative output then :

```
ps -C gedit -L -f
```

- total number of threads of the process in the NLWP column.

**Sample code:**

## 4. Listing Child Processes

see the spawned child processes rather than the threads.

**Command:** `ps -e -H`

- H - for heirarchy
- we can filter by session id (SID) only and not by pid or process name.  
For that: `ps -g sid -H`

or

`ps -s sid -H`

**Sample code:**

```
vboxuser@ubuntu:~$ ps -e -H | head -5
  PID TTY          TIME CMD
    2 ?        00:00:00 kthreadd
    3 ?        00:00:00  pool_workqueue_release
    4 ?        00:00:00  kworker/R-rcu_gp
    5 ?        00:00:00  kworker/R-sync_wq
vboxuser@ubuntu:~$ ps -e -H -o pid,cmd | head -5
  PID CMD
    2 [kthreadd]
    3 [pool_workqueue_release]
    4 [kworker/R-rcu_gp]
    5 [kworker/R-sync_wq]
vboxuser@ubuntu:~$ ps -C bash -o sid
  SID
 2706
vboxuser@ubuntu:~$ ps -g 2706 -H | head -5
  PID TTY          TIME CMD
2706 pts/0    00:00:00 bash
2839 pts/0    00:00:00  ps
2840 pts/0    00:00:00  head
vboxuser@ubuntu:~$ █
```

**How to see my session id?**

`ps -C gedit -o sid`

Note:

session id is equal to the process id that started the session — also called the session leader.

## 5. Controlling the Output

control which columns are printed with the help of the -o flag ( order oriented).

```
ps -C gedit -o pid,tty
```

**Sample code:**

```
sr/share/:/var/lib/snapd/desktop
vboxuser@ubuntu:~$ ps -e -o pid,cmd | head -6
  PID CMD
    1 /sbin/init splash
    2 [kthreadd]
    3 [pool_workqueue_release]
    4 [kworker/R-rcu_gp]
    5 [kworker/R-sync_wq]
vboxuser@ubuntu:~$
```

## 6. Process Tree

**Command:**

```
pstree -p
```

**Explanation:**

Shows parent-child process relationships.

- [] denote identical branches.
- { } denote child thread.

**Example Output:**

```
systemd(1)---NetworkManager(778)
  |-sshd(895)---sshd(1023)---bash(1024)---pstree(1101)
  |  `--mysqld(2001)
  `--python3(1234)
```

**Sample code:**

```

└─wpa_supplicant(863)
vboxuser@ubuntu:~$ pstree -p | head -5
systemd(1)-+-ModemManager(936)-+--{ModemManager}(965)
|                               | -{ModemManager}(974)
|                               | -{ModemManager}(977)
| -NetworkManager(858)-+--{NetworkManager}(935)
|                               | -{NetworkManager}(938)
vboxuser@ubuntu:~$ █

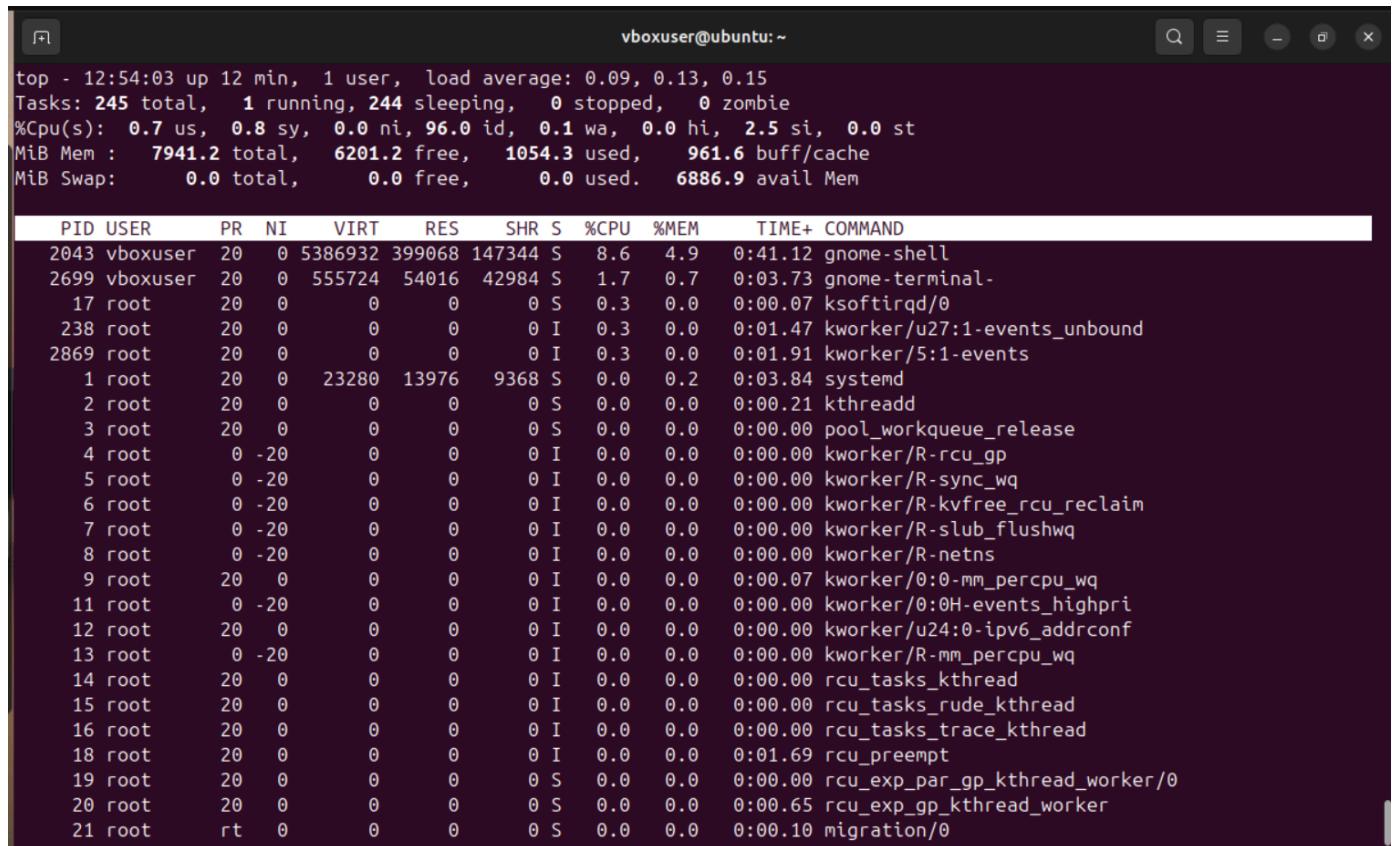
```

## 7. Real-Time Monitoring

### Command:

top

### Sample code:



The screenshot shows the 'top' command running in a terminal window on an Ubuntu system. The title bar says 'vboxuser@ubuntu:~'. The output provides real-time system statistics and a list of processes.

```

top - 12:54:03 up 12 min, 1 user, load average: 0.09, 0.13, 0.15
Tasks: 245 total, 1 running, 244 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.7 us, 0.8 sy, 0.0 ni, 96.0 id, 0.1 wa, 0.0 hi, 2.5 si, 0.0 st
MiB Mem : 7941.2 total, 6201.2 free, 1054.3 used, 961.6 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used. 6886.9 avail Mem

      PID USER      PR  NI    VIRT    RES    SHR S %CPU %MEM     TIME+ COMMAND
 2043 vboxuser  20   0 5386932 399068 147344 S  8.6  4.9  0:41.12 gnome-shell
 2699 vboxuser  20   0 555724 54016 42984 S  1.7  0.7  0:03.73 gnome-terminal-
 17 root      20   0      0      0      0 S  0.3  0.0  0:00.07 ksoftirqd/0
 238 root      20   0      0      0      0 I  0.3  0.0  0:01.47 kworker/u27:1-events_unbound
 2869 root      20   0      0      0      0 I  0.3  0.0  0:01.91 kworker/5:1-events
  1 root      20   0 23280 13976  9368 S  0.0  0.2  0:03.84 systemd
  2 root      20   0      0      0      0 S  0.0  0.0  0:00.21 kthreadd
  3 root      20   0      0      0      0 S  0.0  0.0  0:00.00 pool_workqueue_release
  4 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/R-rcu_gp
  5 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/R-sync_wq
  6 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/R-kvfree_rcu_reclaim
  7 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/R-slub_flushwq
  8 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/R-netns
  9 root      20   0      0      0      0 I  0.0  0.0  0:00.07 kworker/0:0-mm_percpu_wq
 11 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/0:0H-events_highpri
 12 root      20   0      0      0      0 I  0.0  0.0  0:00.00 kworker/u24:0-ipv6_addrconf
 13 root      0 -20      0      0      0 I  0.0  0.0  0:00.00 kworker/R-mm_percpu_wq
 14 root      20   0      0      0      0 I  0.0  0.0  0:00.00 rcu_tasks_kthread
 15 root      20   0      0      0      0 I  0.0  0.0  0:00.00 rcu_tasks_rude_kthread
 16 root      20   0      0      0      0 I  0.0  0.0  0:00.00 rcu_tasks_trace_kthread
 18 root      20   0      0      0      0 I  0.0  0.0  0:01.69 rcu_preempt
 19 root      20   0      0      0      0 S  0.0  0.0  0:00.00 rcu_exp_par_gp_kthread_worker/0
 20 root      20   0      0      0      0 S  0.0  0.0  0:00.65 rcu_exp_gp_kthread_worker
 21 root      rt   0      0      0      0 S  0.0  0.0  0:00.10 migration/0

```

### Explanation:

Displays CPU, memory, and process usage in real-time.

- Press **q** to quit.

### Example Output (partial):

```

top - 10:20:51 up 2 days, 3:12, 2 users, load average: 0.22, 0.33, 0.45
Tasks: 197 total, 1 running, 196 sleeping, 0 stopped, 0 zombie

```

```
%Cpu(s): 12.3 us, 5.4 sy, 0.0 ni, 80.1 id, 2.2 wa, 0.0 hi, 0.0 si, 0.0 st  
KiB Mem : 8045632 total, 3564980 free, 1876324 used, 2604328 buff/cache
```

## 8. Adjust Process Priority

The `nice` command is used to start a process with a specified nice value, while the `renice` command is used to alter priority of the running process.

Syntax:

```
sudo nice -n <nice_value> <command_name>
```

**Value of nice:** -20(highest) to +20(lowest).

**Default:** 0.

-n:

Set **the nice value for a new process**.

-p:

Modify **the nice value for a process using its PID**.

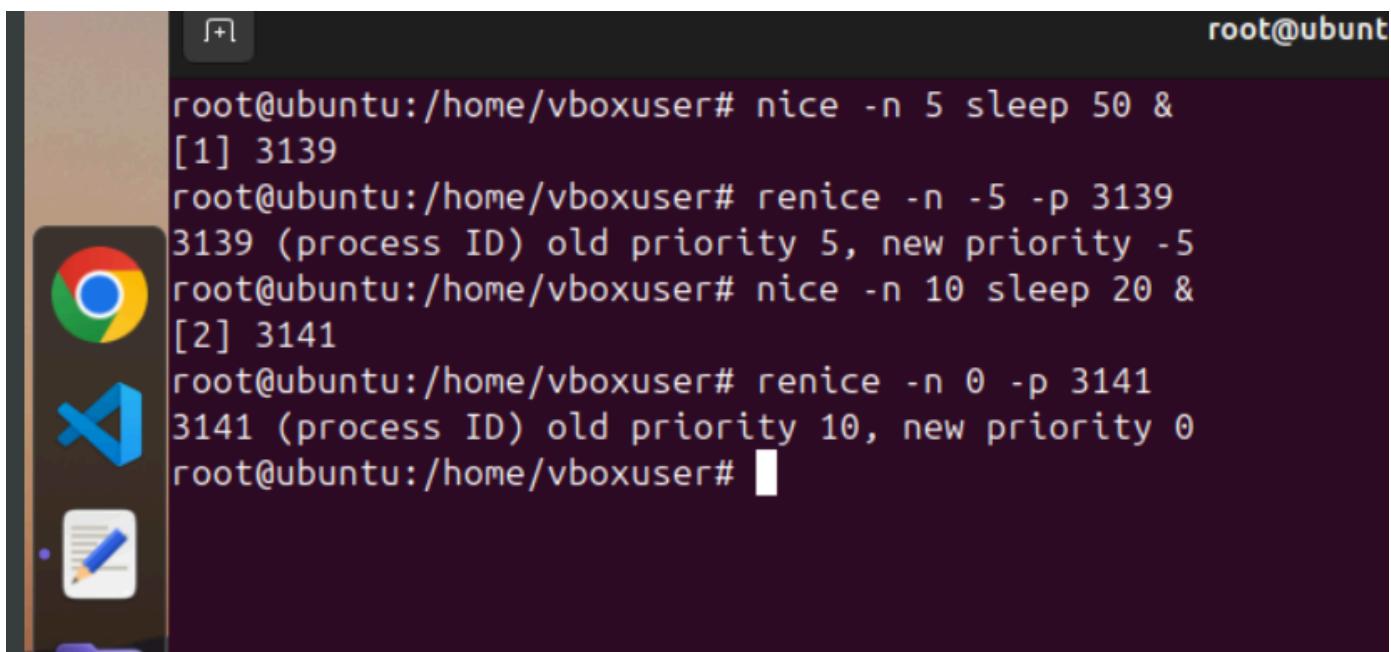
**Start a process with low priority:**

```
sudo nice -n 10 sleep 300 &
```

**Output:**

```
[1] 3050
```

**Sample code:**



A screenshot of a terminal window titled "root@ubuntu". The terminal shows the following session:

```
root@ubuntu:/home/vboxuser# nice -n 5 sleep 50 &
[1] 3139
root@ubuntu:/home/vboxuser# renice -n -5 -p 3139
3139 (process ID) old priority 5, new priority -5
root@ubuntu:/home/vboxuser# nice -n 10 sleep 20 &
[2] 3141
root@ubuntu:/home/vboxuser# renice -n 0 -p 3141
3141 (process ID) old priority 10, new priority 0
root@ubuntu:/home/vboxuser#
```

The terminal window has a dark theme and shows icons for Google Chrome, Microsoft Word, and Microsoft Excel in the background.

**Change priority of `running` process:**

```
renice -n -5 -p 3050
```

Note: You can only use renice when the process is running and hasn't ended.

**Output:**

```
3050 (process ID) old priority 10, new priority -5
```

**Note:** Lower nice values (negative) mean **higher** priority. Requires appropriate privileges.

## 9. CPU Affinity (Bind Process to CPU Core)

**Check CPU affinity:**

```
taskset -cp 3050
```

Here, if you use -p instead of -cp then it just shows 3f (in my system)

**Output:**

```
pid 3050's current affinity list: 0-3
```

**Restrict to core 1 only:**

```
taskset -cp 1 3050
```

**Output:**

```
pid 3050's current affinity list: 1
```

**Sample code:**

```
vboxuser@ubuntu:~$ taskset -p 3270
pid 3270's current affinity mask: 3f
vboxuser@ubuntu:~$ taskset -cp 3270
pid 3270's current affinity list: 0-5
vboxuser@ubuntu:~$ taskset -cp 3271
taskset: failed to get pid 3271's affinity: No such process
vboxuser@ubuntu:~$ taskset -cp 3257
pid 3257's current affinity list: 0
vboxuser@ubuntu:~$
```

## 10. I/O Scheduling Priority

**Command:**

Syntax:

```
ionice -c scheduling_class -n priority_nice_value command
```

**Note:** `ionice` affects Disk I/O not CPU.

Number of the scheduling classes:

- 0 for none
- 1 for real-time
- 2 for best-exertion
- 3 for inactive.

```
ionice -c 3 -p 3050
```

#### Output:

```
successfully set pid 3050's IO scheduling class to idle
```

- **Class 3 (idle)** → Process only gets I/O when system is idle.

#### Sample code:

```
vboxuser@ubuntu:~$ ionice -c 3 bash
vboxuser@ubuntu:~$ ps -e | tail -5
 2988 pts/0    00:00:00 bash
 3000 ?        00:00:00 update-notifier
 3027 pts/0    00:00:00 bash
 3071 pts/0    00:00:00 ps
 3072 pts/0    00:00:00 tail
vboxuser@ubuntu:~$ ionice -c 3 -p 2988
vboxuser@ubuntu:~$ █
```

---

## 11. File Descriptors Used by a Process

#### Command:

**lsof**: The command itself, used to list open files.

**Note:** This command lists out all the files that are opened by any process in the system.

```
lsof -p 3050 | head -5
```

#### Note:

**head -5** shows the first 5 lines of the output incl the column heads.

#### Example Output:

COMMAND	PID	USER	FD	TYPE	DEVICE	SIZE/OFF	NODE	NAME
sleep	3050	vibhu	cwd	DIR	253,0	4096	131073	/home/vibhu
sleep	3050	vibhu	rtd	DIR	253,0	4096	2	/
sleep	3050	vibhu	txt	REG	253,0	17520	133580	/usr/bin/sleep

### Sample code:

```
vboxuser@ubuntu:~$ lsof -p 2988
COMMAND  PID    USER   FD   TYPE DEVICE SIZE/OFF NODE NAME
bash    2988  vboxuser cwd    DIR    8,2      4096  1179650 /home/vboxuser
bash    2988  vboxuser rtd    DIR    8,2      4096       2 /
bash    2988  vboxuser txt    REG    8,2  1446024  1311289 /usr/bin/bash
bash    2988  vboxuser mem    REG    8,2  5719296  1368750 /usr/lib/locale/locale
-archive
bash    2988  vboxuser mem    REG    8,2  2125328  1313198 /usr/lib/x86_64-linux-
gnu/libc.so.6
bash    2988  vboxuser mem    REG    8,2  208328  1322742 /usr/lib/x86_64-linux-
gnu/libtinfo.so.6.4
bash    2988  vboxuser mem    REG    8,2  27028  1313187 /usr/lib/x86_64-linux-
gnu/gconv/gconv-modules.cache
bash    2988  vboxuser mem    REG    8,2  236616  1313195 /usr/lib/x86_64-linux-
gnu/ld-linux-x86-64.so.2
bash    2988  vboxuser  0u    CHR  136,0      0t0       3 /dev/pts/0
bash    2988  vboxuser  1u    CHR  136,0      0t0       3 /dev/pts/0
bash    2988  vboxuser  2u    CHR  136,0      0t0       3 /dev/pts/0
bash    2988  vboxuser 255u   CHR  136,0      0t0       3 /dev/pts/0
vboxuser@ubuntu:~$
```

## 12. Working with Strace Process Monitoring Tool

### Command:

```
strace -p 3050
```

### Example Output:

```
strace: Process 3050 attached
restart_syscall(<... resuming interrupted nanosleep ...>) = 0
nanosleep({tv_sec=300, tv_nsec=0}, 0x7ffd4a60d8b0) = ? ERESTART_RESTARTBLOCK
(Interrupted by signal)
```

Press **Ctrl+C** to detach from `strace`.

### Sample code:

## 13. Find Process Using a Port

## Command:

```
sudo fuser -n tcp 8080  
#used to identify which process is using TCP port 8080
```

## Example Output:

8080/tcp: 4321

PID 4321 is using port 8080.

NOTE: it will output null if no process is running in that port.

## Sample code:

```
vboxuser@ubuntu:~$ sudo fuser -n tcp 1091
vboxuser@ubuntu:~$ sudo fuser -n tcp 0
0/tcp:                                570  1186
vboxuser@ubuntu:~$
```

## 14. Monitoring processes

**Command:**

```
pidstat -p 3050 2 3
```

- pidstat: A performance monitoring tool from the sysstat package.(not in my device ##)
- -p 3050: Monitor only the process with PID 3050.
- 2: Interval in seconds between each report.
- 3: Number of reports to generate.

**Example Output:**

Linux 5.15.0 (ubuntu)		09/25/25		_x86_64_		(4 CPU)	
12:30:20	UID	PID	%usr	%system	%CPU	CPU	Command
12:30:22	1000	3050	0.00	0.00	0.00	1	sleep
12:30:24	1000	3050	0.00	0.00	0.00	1	sleep
12:30:26	1000	3050	0.00	0.00	0.00	1	sleep

**Tip:** `pidstat` is part of the `sysstat` package on most distros.

**Sample code:**

```
vboxuser@ubuntu:~$ pidstat -p 3113 3 3
Linux 6.14.0-27-generic (ubuntu)          09/26/2025      _x86_64_      (6 CPU)

02:56:24 PM  UID    PID  %usr %system  %guest  %wait   %CPU   CPU  Command
02:56:27 PM  1000  3113  0.00  0.00    0.00    0.00   0.00   0.00   2  bash
02:56:30 PM  1000  3113  0.00  0.00    0.00    0.00   0.00   0.00   2  bash
02:56:33 PM  1000  3113  0.00  0.00    0.00    0.00   0.00   0.00   2  bash
Average:    1000  3113  0.00  0.00    0.00    0.00   0.00   0.00   -  bash
vboxuser@ubuntu:~$
```

## 15. Control Groups for Resource Limits

**cgroups**

- a Linux kernel feature that limits, accounts for, and isolates the resource usage (CPU, memory, disk I/O, network, etc.) of a collection of processes.

Q 1: We will create a disk-controlled group so that we may run any process with a finite amount of disk read/writes available. i.e. we want to throttle reads and writes done by a process or a group of

processes.

### Create a new cgroup:

```
sudo cgcreate -g cpu,memory:/testgroup
```

### Limit CPU and Memory:

```
echo 50000 | sudo tee /sys/fs/cgroup/cpu/testgroup/cpu.cfs_quota_us
```

```
echo 100M | sudo tee /sys/fs/cgroup/memory/testgroup/memory.limit_in_bytes
```

### Add a process (PID 3050) to cgroup:

```
echo 3050 | sudo tee /sys/fs/cgroup/cpu/testgroup/cgroup.procs
```

**Note:** Paths and controllers differ between **cgroups v1** and **v2**. Adjust for your system.

### Sample code:

```
root@ubuntu:/# stat -fc %T /sys/fs/cgroup/
cgroup2fs
root@ubuntu:/# cd /sys/fs/cgroup
root@ubuntu:/sys/fs/cgroup# echo +cpu +memory | sudo tee cgroup.subtree_control
+cpu +memory
root@ubuntu:/sys/fs/cgroup# sudo mkdir firstgrp
root@ubuntu:/sys/fs/cgroup# ls -l firstgrp
total 0
-r--r--r-- 1 root root 0 Sep 26 15:17 cgroup.controllers
-r--r--r-- 1 root root 0 Sep 26 15:17 cgroup.events
-rw-r--r-- 1 root root 0 Sep 26 15:17 cgroup.freeze
--w----- 1 root root 0 Sep 26 15:17 cgroup.kill
-rw-r--r-- 1 root root 0 Sep 26 15:17 cgroup.max.depth
-rw-r--r-- 1 root root 0 Sep 26 15:17 cgroup.max.descendants
-rw-r--r-- 1 root root 0 Sep 26 15:17 cgroup.pressure
-rw-r--r-- 1 root root 0 Sep 26 15:17 cgroup.procs
-r--r--r-- 1 root root 0 Sep 26 15:17 cgroup.stat
-rw-r--r-- 1 root root 0 Sep 26 15:17 cgroup.subtree_control
-rw-r--r-- 1 root root 0 Sep 26 15:17 cgroup.threads
-rw-r--r-- 1 root root 0 Sep 26 15:17 cgroup.type
-rw-r--r-- 1 root root 0 Sep 26 15:17 cpu.idle
-rw-r--r-- 1 root root 0 Sep 26 15:17 cpu.max
-rw-r--r-- 1 root root 0 Sep 26 15:17 cpu.max.burst
-rw-r--r-- 1 root root 0 Sep 26 15:17 cpu.pressure
-r--r--r-- 1 root root 0 Sep 26 15:17 cpu.stat
-r--r--r-- 1 root root 0 Sep 26 15:17 cpu.stat.local
-rw-r--r-- 1 root root 0 Sep 26 15:17 cpu.uclamp.max
-rw-r--r-- 1 root root 0 Sep 26 15:17 cpu.uclamp.min
-rw-r--r-- 1 root root 0 Sep 26 15:17 cpu.weight
-rw-r--r-- 1 root root 0 Sep 26 15:17 cpu.weight.nice
-rw-r--r-- 1 root root 0 Sep 26 15:17 io.pressure
-r--r--r-- 1 root root 0 Sep 26 15:17 memory.current
-r--r--r-- 1 root root 0 Sep 26 15:17 memory.events
-r--r--r-- 1 root root 0 Sep 26 15:17 memory.events.local
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.high
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.low
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.max
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.min
-r--r--r-- 1 root root 0 Sep 26 15:17 memory.numa_stat
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.oom.group
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.peak
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.pressure
--w----- 1 root root 0 Sep 26 15:17 memory.reclaim
```

```
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.oom.group
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.peak
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.pressure
--w----- 1 root root 0 Sep 26 15:17 memory.reclaim
-r--r--r-- 1 root root 0 Sep 26 15:17 memory.stat
-r--r--r-- 1 root root 0 Sep 26 15:17 memory.swap.current
-r--r--r-- 1 root root 0 Sep 26 15:17 memory.swap.events
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.swap.high
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.swap.max
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.swap.peak
-r--r--r-- 1 root root 0 Sep 26 15:17 memory.zswap.current
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.zswap.max
-rw-r--r-- 1 root root 0 Sep 26 15:17 memory.zswap.writeback
-r--r--r-- 1 root root 0 Sep 26 15:17 pids.current
-r--r--r-- 1 root root 0 Sep 26 15:17 pids.events
-r--r--r-- 1 root root 0 Sep 26 15:17 pids.events.local
-rw-r--r-- 1 root root 0 Sep 26 15:17 pids.max
-r--r--r-- 1 root root 0 Sep 26 15:17 pids.peak
root@ubuntu:/sys/fs/cgroup# echo $$ | sudo tee /firstgrp/cgroup.procs
tee: /firstgrp/cgroup.procs: No such file or directory
3149
root@ubuntu:/sys/fs/cgroup# cd firstgrp
root@ubuntu:/sys/fs/cgroup/firstgrp# echo $$ | sudo tee cgroup.procs
3149
root@ubuntu:/sys/fs/cgroup/firstgrp# echo "50000 100000" | sudo tee cpu.max
50000 100000
root@ubuntu:/sys/fs/cgroup/firstgrp# echo 104857600 | sudo tee memory.max
104857600
root@ubuntu:/sys/fs/cgroup/firstgrp# cat cpu.stat
usage_usec 185642
user_usec 59247
system_usec 126394
nice_usec 0
core_sched.force_idle_usec 0
nr_periods 16
nr_throttled 0
throttled_usec 0
nr_bursts 0
burst_usec 0
root@ubuntu:/sys/fs/cgroup/firstgrp# cat memory.current
675840
root@ubuntu:/sys/fs/cgroup/firstgrp# █
```

## 16. Alternatives to nice / renice

### 1. `chrt` command

- command-line utility used to set or retrieve the real-time scheduling policy and priority of a process in Linux.

## ⌚ Linux Scheduling Policies: `SCHED_FIFO`, `SCHED_RR`, `SCHED_OTHER`

Policy	Type	Priority Range	Time Slice	Preemption	Use Case
<code>SCHED_FIFO</code>	Real-time	1–99	None	Yes	Deterministic tasks needing strict order
<code>SCHED_RR</code>	Real-time	1–99	Yes	Yes	Fair time-sharing among real-time tasks
<code>SCHED_OTHER</code>	Normal (default)	0	Yes	Yes	General-purpose user processes

### Notes:

- `SCHED_FIFO`: Runs until blocked or preempted. No time slice.  
🚩 FLAG: `-f`
- `SCHED_RR`: Round-robin with time slices. Fair among equal-priority tasks.  
🚩 FLAG : `-r`
- `SCHED_OTHER`: Default Linux scheduler (CFS). Used by most user-space processes.  
🚩 FLAG: `-o`

```
vboxuser@ubuntu:~$ sudo -i
[sudo] password for vboxuser:
root@ubuntu:~# pkill -f "sleep 300"
root@ubuntu:~# chrt -r 40 sleep 300
^C
root@ubuntu:~# pkill -f "sleep 300"
root@ubuntu:~# chrt -r 40 sleep 300 &
[1] 3943
root@ubuntu:~# pgrep -f "sleep 300"
3943
root@ubuntu:~# chrt -p 3943
pid 3943's current scheduling policy: SCHED_RR
pid 3943's current scheduling priority: 40
root@ubuntu:~#
```

### 2. ionice (I/O Priority Control)

Refer to: [I/O Scheduling Priority](#)

### 3. taskset (CPU Affinity)

Refer to : [CPU Affinity \(Bind Process to CPU Core\)](#).

### 4. Control Groups (cgroups)

Refer to : [Control Groups for Resource Limits](#)

### 5. systemd-run

Code:

```
systemd-run --scope -p CPUWeight=200 stress --cpu 4
```

Output:

launching the stress tool inside a transient systemd scope unit with a CPU weight of 200.

Explanation:

- `systemd-run`: Creates and starts a transient systemd unit.
- `--scope`: Runs the command in a scope unit, which is suitable for processes that are not started by systemd itself (like interactive commands).
- `-p CPUWeight=200`: Sets the CPU weight for the scope unit. This is part of systemd's CPU scheduling via cgroups. The default weight is 100; higher values give more CPU time relative to other units.
- `stress --cpu 4`: Runs the stress tool to spawn 4 CPU workers, each spinning in a tight loop to simulate CPU load.

```
vboxuser@ubuntu:~$ sudo apt install stress
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following package was automatically installed and is no longer required:
  libllvm19
Use 'sudo apt autoremove' to remove it.
The following NEW packages will be installed:
  stress
0 upgraded, 1 newly installed, 0 to remove and 36 not upgraded.
Need to get 18.1 kB of archives.
After this operation, 52.2 kB of additional disk space will be used.
Get:1 http://in.archive.ubuntu.com/ubuntu noble/universe amd64 stress amd64 1.0.7-1 [18.1 kB]
Fetched 18.1 kB in 1s (29.0 kB/s)
Selecting previously unselected package stress.
(Reading database ... 202873 files and directories currently installed.)
Preparing to unpack .../stress_1.0.7-1_amd64.deb ...
Unpacking stress (1.0.7-1) ...
Setting up stress (1.0.7-1) ...
Processing triggers for man-db (2.12.0-4build2) ...
vboxuser@ubuntu:~$ systemd-run --scope -p CPUWeight=200 stress --cpu 4
Running as unit: run-u111.scope; invocation ID: ceea114cb00f48d1909ed1d31ab24431
stress: info: [3641] dispatching hogs: 4 cpu, 0 io, 0 vm, 0 hdd
^C
vboxuser@ubuntu:~$ top

top - 06:11:52 up 15 min,  1 user,  load average: 0.86, 0.82, 0.56
Tasks: 248 total,   1 running, 247 sleeping,   0 stopped,   0 zombie
CPU states:  0%us,  0%sy,  0%ni, 99%id,  0%wa,  0%hi,  0%si,  0%st

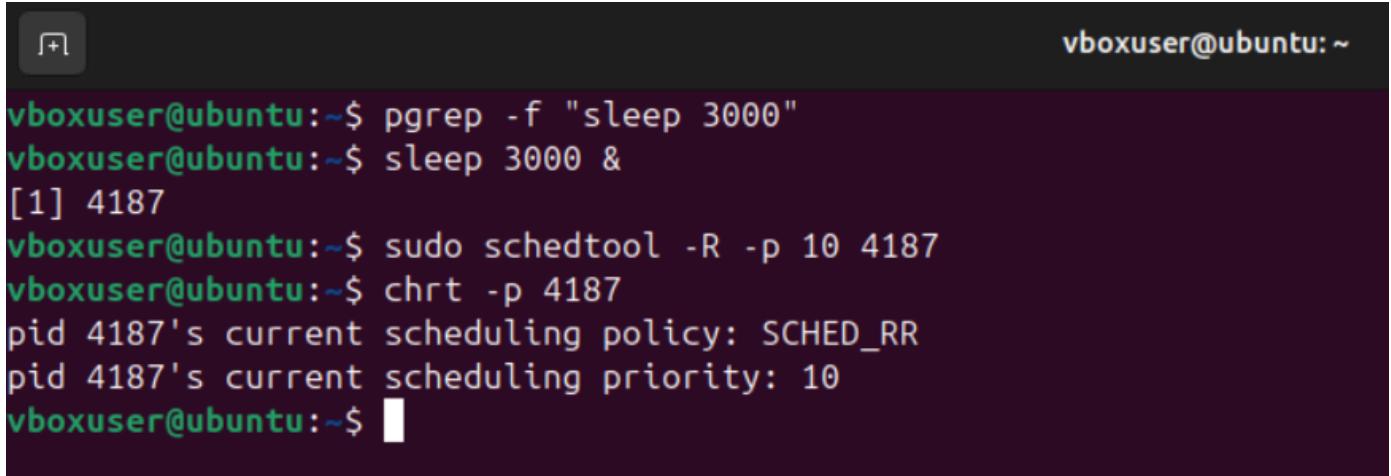
```

### 7. schedtool

Syntax:

```
sudo schedtool -R -p 10 <pid>
```

Sample code:



A screenshot of a terminal window titled "vboxuser@ubuntu: ~". The terminal shows the following sequence of commands:

```
vboxuser@ubuntu:~$ pgrep -f "sleep 3000"
vboxuser@ubuntu:~$ sleep 3000 &
[1] 4187
vboxuser@ubuntu:~$ sudo schedtool -R -p 10 4187
vboxuser@ubuntu:~$ chrt -p 4187
pid 4187's current scheduling policy: SCHED_RR
pid 4187's current scheduling priority: 10
vboxuser@ubuntu:~$
```

## 17. ❤ Cheat Sheet

Command	Purpose
<code>ps -C processname</code>	Filter processes by name
<code>ps -p PID</code>	Show details for specific PID
<code>ps -u username</code>	List processes by user
<code>ps -C gedit -L</code>	List threads of a process
<code>ps -e -H</code>	Show child processes in hierarchy
<code>ps -C gedit -o pid,tty</code>	Customize output columns
<code>strace -p PID</code>	Monitor system calls of a process
<code>sudo fuser -n tcp PORT</code>	Find process using a TCP port
<code>pidstat -p PID interval count</code>	Monitor CPU usage over time
<code>cgroup -g cpu,memory:/group</code>	Create a control group
<code>echo PID   tee /sys/fs/cgroup/...</code>	Add process to cgroup
<code>systemd-run --scope -p CPUWeight=200 command</code>	Run process with CPU weight control
<code>schedtool -R -p N PID</code>	Set real-time scheduling policy